

Cloud Composer v2 - Michaël Bettan

Definition: Managed Apache Airflow workflow orchestration service. Apache Airflow is an open source platform to orchestrate data pipelines programmatically in Python.

Availability:
Regional

Use Cases:

- Orchestration of data ingestion, including hybrid use cases
- Data processing
- Orchestration of proprietary API endpoints jobs
- Custom code tasks run directly in Airflow
- ML / AI pipelines

Billing:

- Web, Database **core hours**
- Database SQL **core hours**
- Web and Data **storage**
- **Network egress** for all
- Workers and Scheduler **GKE nodes**
- **GCS Bucket** (Dags)
- **Cloud Monitoring** logs usage

Key Capabilities

- Author, Schedule and Monitoring within Airflow
- Composer v2 with Python 3 is the current standard
- OSS making a great fit hybrid & multi-cloud architecture

Airflow DAGs

- **DAG:** Directed Acyclic Graphs. Acyclic = not a circle
- **Workflow pipeline** and **DAG** are interchangeable
- **DAGs:** executed workflows
 - Collection of tasks with dependencies/relationships
 - Stored in Cloud Storage
 - Supports custom plugins: operators, hooks and interfaces
 - Python dependencies (modules)
- **Example:** Export BigQuery data → Move GCS folder → Cleanse data → Ingest in another system (e.g., Snowflake) → Send Slack Message
- DAGs folder is a Cloud Storage bucket where you will load your pipeline code

Airflow Operators

- Core building blocks of DAGs. **Airflow operators** invoke the **TASKS** you want to complete. Operators are usually atomic in a task (one operator per task).
- Represent individual tasks within a workflow.
- Types:
 - **DummyOperator:** visual clarity(start/end points).
 - **BashOperator:** Executes bash commands.
 - **PythonOperator:** Executes Python code.
 - **EmailOperator:** Sends emails.
 - **GoogleCloudStorageToBigQueryOperator:** Transfers data from Cloud Storage to BigQuery.
 - **BigQueryToCloudStorageOperator:** Exports data from BigQuery to GCS. Pay attention to the different export formats (Avro, Parquet) and compression options (Snappy, Deflate) for cost optimization.
 - **GoogleCloudStorageToGoogleCloudStorageOperator** : Copies files between GCS buckets.
 - **Many others:** for interacting with other services (e.g., DataProc, Dataflow, Pub/Sub).

Architecture

- **Environment** is the instance construct, deployed automatically with several components:
 - Customer project → GKE, Cloud Storage
 - Google-managed tenant project → Cloud SQL, Identity-Aware-Proxy, App Engine Flexible
- **Customer project** is where you create your environments. You can create more than one environment in a single customer project.
- **Google-managed tenant project** provides unified access control and an additional layer of data security for your environment. Each environment has its own tenant project.
- **Airflow web server** on App Eng Flex is the Airflow UI
- **Airflow database** on Cloud SQL: stores the metadata
- **Environment's bucket** on GCS to store DAGs, plugins, data dependencies, and Airflow logs.
- **Airflow scheduler** controls the scheduling of DAG runs and individual tasks from DAGs. Distribute tasks to Airflow workers by using a Redis queue, as a GKE deployments.
- **Airflow workers** execute individual tasks from DAGs by taking them from the Redis queue. Airflow workers run as GKE deployments.
- **Redis queue** holds a queue of individual tasks from your DAGs. Airflow schedulers fill the queue; Airflow workers take their tasks from it. Redis queue runs as a GKE StatefulSet application, so that messages persist across container restarts.
- For each Private IP environment, Cloud Composer creates one **VPC peering connection** for the tenant project network

Logging & Monitoring

- 2 types of logs: operational (system) and task
- Operational logs (Scheduler) -- Logs Viewers in the console (from Cloud Monitoring)
- Each DAG task has an associated log via Airflow UI, Logs folder (GCS)
- Monitoring natively from the Environment UI directly

Security

- Composer Administrator: full control of resources
- Environment and Storage Object Administrator
- Environment User and Storage Object Viewer
- Composer User

Failure and Retry Mechanisms

Airflow provides robust mechanisms to handle task failures.:

- Users can define **retry logic within their DAGs**, specifying the **number of retry attempts** and the **time interval between retries**.
- On failure, Airflow can **trigger alerts and store logs** for debugging, enabling automated recovery and efficient troubleshooting
- Specific operators may also offer **customized failure handling options**.
- Tasks in a DAG can fail for various reasons. The **on_failure_callback parameter** is designed to handle situations where a task fails, allowing custom logic (such as sending notifications) to be executed when the task does not succeed.

ETL: Push vs. Pull patterns

- **Push-based ETL with Cloud Functions and Airflow:** Data ingestion is initiated by an external event (e.g., a new file in Cloud Storage, a Pub/Sub message). This event triggers a Cloud Function, which in turn triggers the execution of a specific Airflow DAG responsible for the ETL process. This approach is **reactive and ideal for near real-time data processing** where the ETL process is dependent on the arrival of new data. Cloud Function acts as a *lightweight intermediary*, decoupling the event source from the more resource-intensive Airflow DAG execution. Efficient for handling variable data volumes and maintaining responsiveness.
- **Pull-based ETL with Scheduled Airflow DAGs:** initiated on a **predefined schedule** (e.g., cron expression) directly within Airflow. Airflow Scheduler autonomously checks for scheduled DAG runs and executes them at the designated times. This is a proactive approach best suited for **batch processing** of data at **regular intervals** where data volume is more predictable. It's simpler to manage than push-based systems but might be less responsive to immediate data changes.